

Detection
 $p(\text{buggy}) = 0.98$

Run: BERT / MRPC
embedding_swap / L1

Family

Embedding
 $p = 0.965$

refines

Root cause

embedding_swap
 $p = 0.994$

Top subsystems

representation
 $A = 2.52$

runtime
 $A = 1.23$

positional
 $A = 1.18$

output_logits
 $A = 1.02$

Top evidence

repr_l0_drift
 $z = 3.07$

peak_mem_alloc
 $z = 1.65$

repr_l1_drift
 $z = 1.97$

logit_margin
 $z = 2.04$

Fix checklist

1. Verify embedding table shape: (vocab_size, d_model)
2. Check embedding norm distribution across vocabulary
3. Verify num_token_types matches task (2 for sentence pairs)
4. If norm near-zero: initialization or zeroing bug

*Structured
Diagnostic
Report*